

Practical ML Tutorial: Part I


—

SEEMAPLD2023, George Williams

Agenda

Part I

- AI Trends
- ML Basics
- Survival Analysis
- Hands-On Programming



Part II

- AI Hardware
- PyTorch Basics
- Computer Vision
- Hands-On Programming



AI Trends

Have You Been Living Under A
Rock Recently?

- ChatGPT
 - Stable Diffusion
 - Deep Fakes
 - Alpha Fold
 - Foundation Models
-



OpenAI

ChatGPT Introduced





**STABLE
DIFFUSION**



DALLE 2



MIDJOURNEY


nature



DeepMind

Machine Learning Basics





Machine Learning Basics




Data

Machine Learning Basics



Data




Model

Machine Learning Basics



Data




Model



Eval

Machine Learning “In The Wild”




Survival Analysis

—

When Will It End?


When Will *It* End?

When Will *It* End?



Customer Churn

When Will *It* End?




Customer Churn



Health Outcomes

When Will *It* End?



Customer Churn




Health Outcomes



Machine Failure


censoring
period^{free}
time
event

Survival^{drops}
follow-up
observed
follow
probability^{interest}
loss^{right}
censoring
observation




$$S(t) = \Pr(T > t)$$

Hazard Function




Survival Function



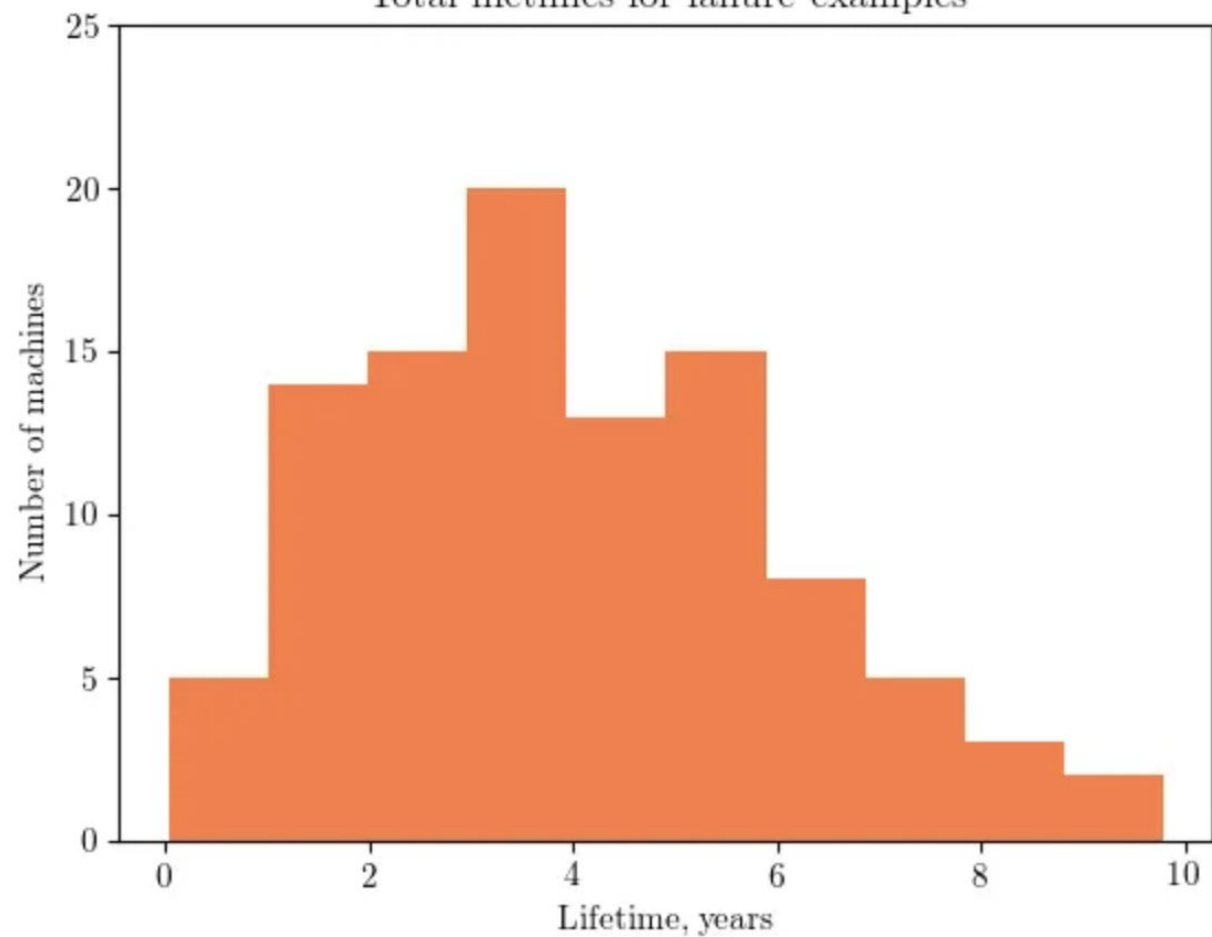
would like to use observed data in order to predict the remaining life for each of a group of machines. To understand how this can be done, it's easiest


Weibull distributions



Shape
Scale

Total lifetimes for failure examples





Machine Learning Approach

Beyond Curve-Fitting

- Handle real-world data
 - Multiple underlying distributions
 - Censored data
 - Best of both worlds: distribution-driven + data-driven
-

Let's Start Coding!

—